hsfpy - A Python Interface to the Hierarchical spline forest C++ Library

Spencer Lyon

Monday 5th August, 2013

A senior capstone project submitted to the faculty of
Brigham Young University
in partial fulfillment of the requirements for the degree of

Bachelor of Science

Dr. Derek Thomas, Advisor

Department of Physics and Astronomy

Brigham Young University

August 2013

ABSTRACT

hsfpy - A Python Interface to the Hierarchical spline forest C++ Library

Spencer Lyon
Department of Physics and Astronomy
Bachelor of Science

I describe the creation of a Python interface to the HSF C++ library. HSF stands for hierarchal spline forests and the C++ library is used to represent surfaces or volumes of arbitrary complexity in terms of hierarchal splines. This library is under active development by BYU faculty in the Physics and Engineering departments. I will defend the choice of using Python as the high-level interface. I will also describe projects that facilitate wrapping compiled languages (like C, C++ or Fortran) in Python. Among them are SWIG, Boost.Python, Cython, and a relatively new project – XDress. XDress blends an expressive typesystem, C/C++ source code parsers, and code generating utilities into an easy to use system for constructing Python wrappers for C or C++ code via Cython.

CONTENTS

`hsfpy` - A Python Interface to the Hierarchical spline forest C++ Library

1 INTRODUCTION

1.1 BACKGROUND

A physicist is interested in discovering and explaining why things are the way they are. This is usually done by making observations, isolating important variables or factors, and building models. In order to use and solve these models physicists need a way to represent them visually and/or in terms of mathematical functions. Especially in physics, these mathematical functions are differential or difference equations with an associated set of boundary conditions.

There are many numerical techniques commonly employed to solve boundary value problems. Among them are finite difference methods (FD), finite element methods (FEM), boundary element methods (BEM), and finite volume methods (FVM). A 2-component approach is taken with each of the techniques:

1) The geometry of the problem is discretized and represented using a mesh.

2) One of the above methods is applied to this mesh and a solution to the model is computed.

Each of the discrete solution techniques mentioned above has its own strengths and weaknesses. FD methods are relatively easy to implement, but are restricted to rectilinear geometry[1]. FEM, BEM, and are all more flexible in how the geometry can be represented, but are typically more difficult to implement. In section 1.1.1, I describe where this additional flexibility comes from.

---

[1]It is actually possible to systems with more complex geometries, but it is difficult and using another method is suggested

1.1.1 SPLINES

Often the mesh used in FEM, BEM, or FVM is defined using a system of splines. A spline is a smooth, piecewise defined polynomial function that is also smooth where the polynomials pieces come together [1]. For numerical methods, the standard type of spline to use to represent the geometry is the non-uniform rational B-spline (NURBS). The definition of a NURBS starts with a non-decreasing and potentially non-uniform vector of knots, which discretize the domain into smaller regions. Polynomial functions are then defined on each of those regions. Next, a set of weights is applied to the basis functions to define rational functions. Finally, a set of coefficients link the rational basis functions to the geometry that is to be described.

Once the geometry has been described (with a grid, NURBS or otherwise), the next step in solving a boundary value problem is to perform the analysis on the discrete system. One inefficiency with standard FEM, BEM, and FVM methods is that the mathematical constructs used to represent the geometry are different than those used to perform the analysis. The analysis items are given many names, such as triangular patches, square patches, tetrahedral patches, or hex patches. Any of these patches is a transformed version of the geometrical representation (i.e. NURBS). This poses at least three issues: 1) It is computationally costly to move from one representation to another, 2) it requires a significant amount of labor to convert to and from the analysis and design geometries, and 3) conversion between geometries introduces some error.

1.1.2 ISOGEOMETRIC ANALYSIS (IGA)

To overcome these issues, a new computational approach called isogeometric analysis (IGA) was introduced in 2005 by Hughes et al. [2]. The main idea behind IGA is to use the exact same basis functions to represent the geometry *and* do the analysis [3]. This simple idea streamlines the interaction between geometric design and rigorous analysis. IGA also provides many other benefits to the design and analysis process. FEM and BEM can use the smooth, high-order basis functions that describe the geometry to perform the computation and analysis, resulting in more accurate results. Also, BEM are usually plagued by geometric error; IGA removes the need to

2

convert between geometries and completely eliminates this error.

1.1.3   HSF

Hierarchal spline forests (HSF) are the focus of current research at BYU and bring additional improvements on top of IGA. Groups of NURBS can be organized into nested, hierarchal structures called spline trees. The spline trees can then be collected as an unstructured, geometrically conforming arrangement called a spline forest. This forest gives IGA a number of enhancements, among which are the following:

- HSF basis functions have compact support and can be made into a partition of unity.

- HSF curves can be made C^∞ between knots and C^{p-k} at knots (p is the degree of spline, k is multiplicity of knot). In this way the user can control the degree of continuity at knot locations.

- Local refinement of basis functions is possible (not generally true of splines).

- Solutions to boundary value problems obtained using HSF curves are both accurate and smooth.

- Geometric structure of governing PDEs can be incorporated directly into the basis (for example $\nabla \cdot \mathbf{B} = 0$ in EM, or $\nabla \cdot \mathbf{v} = 0$ in incompressible flow).

To accompany the theory behind HSF, a C++ library is being developed that implements these concepts (`hsf` will henceforth refer to the C++ library). C++ was an appropriate language choice for the implementation of `hsf` for a number of reasons:

- C++ is a statically typed, compiled programming language. This allows code written in C++ to be executed very fast. For the types of problems IGA and the HSF theory are usually applied to, this is an absolute must.

- C++ is an object oriented programming language. This programming paradigm allows the ideas behind `hsf` (NURBS trees and spline forests, ect.) to be expressed in a very natural way.

- C++ is a mature and has a great foothold in the scientific community. This means that there are many highly optimized and well-tested libraries available for use in `hsf`.

- Some advanced language features, like templates and method, function, or operator overloading, allow the code be general, but still compiled.

On the other hand, there are some shortcomings to choosing C++ as the primary programming language for `hsf`:

- Because C++ is a compiled language, quite a bit of time and effort was spent creating a robust, cross-platform build system for `hsf`.

- C++ is a relatively low-level language. While this does mean it can achieve great performance, it also means that the language is difficult to learn. This can be a barrier to entry for people, especially undergraduate students, who would like to contribute to the development of `hsf`[2].

- Also due to the low-level nature of C++, it tends to be more verbose than other languages. The amount of C++ code required to do a task is often much more than the amount of Matlab or Python code required to do the same thing.

The vision for the `hsf` library is that it will become the most powerful and flexible discretization package for engineering and physics. That it is implemented in C++ gives `hsf` the potential of being very powerful, but might also limit its user base. For that reason, the research group has decided on building an interface between the core C++ library and a higher-level language.

1.2   MOTIVATION

There are many possible options for a high-level interface to `hsf`. Among the most common are Matlab, R, Julia and Python. Each of these languages has its respective strengths. Matlab is among the most popular languages for high-level numerical analysis and computation. R is the standard for open-source statistical programming. Julia couples a dynamic typesystem and advanced multiple dispatch paradigm with a powerful just-in-time compiler to achieve excellent

---

[2]This is apparent in that the main developers of the library are all faculty members.

performance for numerical programming tasks[3]. Python, in contrast, features a complete, state of the art scientific analysis framework build on top of a fully functional programming language.

We decided to use Python to build the interface to C++ for a number of reasons. Python has long had a reputation for being a good "glue" language. The core of the most common implementation of Python, CPython, is actually written in C and thus boasts a native Python-C API. In many ways, the environment most similar to Python is Matlab, but Matlab comes with a hefty price-tag. Python is free, open source, and runs on almost all operating systems. Python is known for its very readable syntax. It is not unreasonable to expect a researcher to be introduced to Python in the morning and be writing meaningful programs by the end of the day[4]. Additionally, in Python it is relatively easy to employ parallel processing. The package mpi4py exposes any system implementation of the message passing interface (MPI) to Python. `hsf` is currently using MPI to implement core algorithms in parallel. Being able to use MPI from Python will help increase the rate of development for parallel `hsf`.

These virtues of Python language all come together into the ideal programming environment for the high-level C++ interface, which we call `hsfpy`. The hope is that a robust and functional implementation of `hsfpy` will assist in the larger goal of `hsf` becoming the go-to package for discretization by lowering the bar of entry. This will allow more researchers to use `hsf` to do their analysis and more students who would like to contribute to the development of `hsf` itself.

The core C++ library for `hsf` is still being actively developed, but is at a very mature state. As of August 2, 2013 there are over 18,000 lines of actual code (excluding blank lines and comments)[5] in the library. This has provided a very stable base upon which the Python interface has been developed.

---

[3]Julia also supports native calls to C/C++ through the `ccall` method

[4]Obviously a mastery of the language will develop over time, but the point of Python being readable and easy to learn stands.

[5]This was determined using the `cloc` utility

2 METHODS

In this section I describe the different approaches that were employed and considered during the creation of `hsfpy`. I will give an overview of the tools that were considered for this project as well as a short usage example for each tool. To maintain consistency and make differences across the methods more apparent, I will use a selection of the code from the `hsf` C++ library. The main components of this example code are a C++ class `HKnotVector`, a function `numClamp`, and a few typedefs, `DoubleVec, IntVec`, and `IntVecVec`. The goal for each example will be to correctly wrap the `HKnotVector` class, as it uses the other components. The source code can be found in Appendix A.

2.1 SWIG

SWIG[6] is an acronym meaning simplified wrapper and interface generator. The following excerpt from the SWIG homepage provides a good explanation of what SWIG is commonly used for:

> SWIG is a software development tool that connects programs written in C and C++ with a variety of high-level programming languages. SWIG is used with different types of target languages including common scripting languages such as Perl, PHP, Python, Tcl and Ruby...SWIG is most commonly used to create high-level interpreted or compiled programming environments, user interfaces, and as a tool for testing and prototyping C/C++ software. SWIG is typically used to parse C/C++ interfaces and generate the 'glue code' required for the above target languages to call into the C/C++ code.

SWIG is a very well-established project; the first version appeared in July 1995 and the most recent version was released in May 2013. Over the years, SWIG has developed in to a very powerful and flexible tool. The best expression of this flexibility is that SWIG officially has at least partial support for nineteen different target languages, whereas other tools that will be discussed in this section are Python specific. A great aspect of this flexibility is that users can run SWIG on the exact same set of files and generate wrappers for different target languages by simply changing a single command line argument.

---

[6]SWIG is free and open source. The source code is hosted at https://github.com/swig/swig and the homepage for the project is http://www.swig.org/.

However, SWIG is not a perfect tool. Due in part to the freedom it gives users to choose amongst multiple output languages, SWIG generates wrapper code that is relatively difficult to customize for a specific target language. Due to its multiple output paradigm, SWIG doesn't fully support all C++ language features for every output language. This is significant issue for this project, because `hsf` heavily uses advanced C++ techniques. Furthermore, in order to use SWIG, a user must supply an additional interface file (commonly with a `.i` suffix) in which the user uses a C-like syntax to describe the desired interface. Finally, the last main drawback I noticed when testing SWIG for `hsfpy` is that the building/compiling phase for SWIG is non-trivial.

2.1.1 SWIG USAGE EXAMPLE

To give an idea of how to use SWIG, I outline how to construct a Python interface to the code contained in Appendix A. The first step is to create a SWIG interface file where the desired wrapper is designed. I will present the wrapper used to expose the class `HKnotVector`, and then explain the key components.

Listing 1: `HKnotVector.i`: SWIG interface for HKnotVector

```
1  %module hsfpy
2
3  %{
4  #include "../HKnotVector.h"
5  %}
6
7  %include "std_vector.i"
8  namespace std {
9      %template(IntVec)  vector<int>;
10     %template(DoubleVec) vector<double>;
11     %template(IntVecVec) vector<vector<int> >;
12 }
13
14 %import "../common.h"
15 %include "../HKnotVector.h"
```

- **Line 1** Declare the name of the module. In large projects the module name allows SWIG to create wrappers that don't have issues with namespace resolution.

- **Lines 3-5** This is a special block that is copied and pasted, with out SWIG parsing, directly into the generated C/C++ portion of the wrapper. If there are things that need to happen for the underlying source to function, but SWIG doesn't need to know about, they go here.

7

- **Lines 7-12** Notice the use of the %include where C++ programmers are used to seeing #include. This is a special SWIG statement that instructs SWIG to access the file "std_vector.i" (included as part of SWIG) and give the interface access to the vector class from within the namespace std. I then then expose the typdefs found in "common.h" as swig templates.

- **Line 14** The SWIG %import directive is used to tell the wrapper that important items live in common.h, but that no wrapper code needs to be generated for that file.

- **Line 15** Finally the SWIG %include directive is used to include the main file HKnotVector.h in the generated wrapper.

Although the interface file is only 15 lines long, it is fairly complex. One thing to note about this interface is that when it is run, the entire HKnotVector class (really everything defined in HKnotVector.h) is wrapped and exposed to the target language. This could pose problems if various types, functions, or class attributes shouldn't be accessed outside of C or C++.

Using this file is a two-step process: 1) Run SWIG on the "HKnotVector.i" and generate the interface, 2) incorporate the generated files into a build system so that they can be imported into Python. This first step is very straightforward and can be accomplished by running the following from the command line:

```
swig -c++ -python HKnotVector.i
```

This command runs SWIG, tells it that the source language is C++, the target language is Python and that the interface file is HKnotVector.i. After running the command two files will be generated HKnotVector_wrap.cxx and hsfpy.py. Together these files files make up the wrapper of HKnotVector.

The next step is to incorporate these files into a build system so that they can be compiled in a way that the system Python can interact with them. The SWIG documentation gives a few possible methods for doing this, but the recommended solution is to let Python handle the compiling. This will ensure that the correct libraries are linked at compile time and that the version of Python directing the compilation will be able to use the compiled extensions. To do

this, a setup.py file must be created. The setup.py file for this example appears below (note that an explanation of key parts of the file is given after the code).

Listing 2: setup.py file for SWIG

```python
#!/usr/bin/env python
"""
setup.py file for building SWIG hsfpy extensions
"""

from distutils.core import setup, Extension

h_knot_vector = Extension('_hsfpy',
                          sources=['./HKnotVector_wrap.cxx']
                          )

setup(name='hsfpy',
      version='0.1',
      author="Spencer Lyon",
      description="""Wrapping hsfpy for python using SWIG""",
      ext_modules=[h_knot_vector],
      py_modules=["hsfpy"],
      )
```

- **Line 6** From the Python `distutils` package, import the `setup` function and the `Extension` class. The `setup` function is the main driving point in this file and will direct the compilation. Objects of type `Extension` hold all the information the `setup` function needs to compile the extensions.

- **Lines 8-10** Describe the `HKnotVector` extension. Notice the first argument given to the `Extension` constructor is `"_hsfpy"`. This argument tells the `setup` function what to name the shared object (or dynamic linking library on Windows) where the compiled wrapper will be placed. Without custom configuration, SWIG requires that this name be a leading underscore followed by the `%module` name defined in the interface file.

- **Lines 12-18** Call the `setup` function to build all the `Extensions` in the `ext_modules` list. This is also where other metadata about the project goes.

The final step in building the interface is to have Python compile the wrappers. This is done on the command line with a single command:

```
    python setup.py build_ext --inplace
```

This command tells the system Python (whatever `python` resolves to on the user's `$PATH`) to build the extensions outlined in `setup.py` inplace, meaning in the current working directory.

In the end, I decided not to use SWIG to create the interface to the entire `hsf` library. The verbose C-like interface files and the need to create a separate interface file for each source file made SWIG more difficult than necessary. In addition, the fact that all code from an exposed C++ source file is wrapped was overkill for this project. Also, not all C++ languages features that appear in `hsf` are supported by SWIG. However, as can be seen from this small exercise, it is a fairly straightforward, if tedious, process to use SWIG to create a Python interface to C++ code.

2.2 BOOST.PYTHON

Boost.Python[7] (henceforth Boost for short) is an alternative to SWIG and is a highly specialized tool for wrapping C++ for Python use. This apparent lack of flexibility has allowed the Boost developers to provide a very natural and complete coverage of the C++ language. Some key C++ features that are supported in boost are

- References and Pointers
- Efficient function overloading
- C++ to Python exception translation (cuts down on `SEGFAULTs`)
- Functions or methods with default and keyword arguments
- Exporting C++ iterators as Python iterators
- Control over Python documentation strings

On the other hand, Boost has significant limitations. First, Boost has a difficult `Bjam` utility for compiling the wrappers. `Bjam` is similar to `make`, but has a difficult and strange syntax. Second, the generated wrapper code is generally very verbose. While this is is probably due to supporting some C++ features that other wrapping tools do not, it has at least two major drawbacks: 1) It takes a long time to compile the wrappers and 2) the Python-side execution is

---

[7]Boost.Python is part of the free peer-reviewed Boost project. Boost can be downloaded from the main projects webstie at `http://www.boost.org/`. The documentation for Boost.Python can be found at `http://www.boost.org/doc/libs/1_54_0/libs/python/doc/index.html`.

typically noticeably slower than the code generated with other tools. Finally, the major drawback and ultimate reason why I did not use Boost for `hsfpy` is that it is very difficult to install. After reading the (sparse) documentation and searching the internet, I still could not get Boost.Python correctly installed and configured on my system. This would be a major roadblock to future users of the Python bindings and would actually detract from the main justification for creating the bindings: lowering the bar to entry for using `hsf` in research. For these reasons, I will not include a usage example for Boost.Python, but because I spent quite a bit of time on it and many people seem to like it, I felt it needed to be addressed in this report.

2.3   CYTHON

The next tool I examined was `Cython` [8]. Instead of being a tool used solely to wrap compiled languages for use in Python, Cython is actually a super-set of the Python language; anything that is valid Python code is also valid Cython code. However, Cython adds a few major improvements:

- Variables, functions, and classes can be given static types. This avoids much of the overhead inherent in a "duck-typed" interpreted language like Python.

- Cython programs can make direct calls to C, C++, and Fortran code. This allows the user to directly mix Python with low-level, high performance compiled code.

Cython accomplishes this by translating the Cython code directly to C or C++, which can then be compiled and loaded into any Python script or session. This means that blocks of code where all objects have been given static types can be written directly in C and therefore achieve almost[9] C-like performance. In addition, the ability to directly call C, C++, or Fortran makes Cython a viable option for wrapping low-level code for use in Python. In Appendix B, I provide a detailed example static typing in Cython.

---

[8]Cython is free and open source. The source code is hosted at https://github.com/cython/cython and the homepage for the project is http://cython.org/.

[9]The almost is necessary because there is small overhead in calling the compiled routines from Python and getting the results back.

2.3.1 CYTHON WRAPPING EXAMPLE

Building on the static typing example in Appendix B, I will now show how to use Cython to wrap the HKnotVector example[10]. In order to use external libraries, you need to tell Cython two things: 1) what file the external components are defined in (usually a header `.h` file) and 2) which parts of that file you would like to access from Cython. For example, instead of calling `from libc.math cimport sqrt`, I could have done the following:

```
cdef extern from "math.h":
    double sqrt(double x)
```

In the first line I started a `cdef extern` block. The syntax is simply `cdef extern from <headerName>:`, where `headerName` is the name of the external file where the desired objects are defined (`"math.h"` for this example). Everything in the indented block following the `:` is part of this `cdef extern` block and contains the external declarations that need to be exposed to Cython.

In a larger project, it is often necessary to create a Cython interface file (with a `.pxd` extension), which does for Cython what a `.h` interface file does for C/C++. This is necessary when you have multiple Cython `.pyx` files that need to access the same external source. The declarations go into a `cdef extern` block in a `.pxd` file. This interface is very similar to the C/C++ interface; often users can copy and paste directly from C to Cython. The actual implementation will go into a file with the same name, but with a `.pyx` extension. This is very similar to `.h` and `.c` files for C. Furthermore, when wrapping a set of C++ classes, people often put the extern definitions in a file named something like `cpp_<headerName>.pxd` and Cython declarations in a file named `<headerName>.pxd`. This is important because it is generally necessary to have one interface file for external declarations (the file named `cpp_`), and another interface file exposing the Cython implementation.

The structure of a Cython wrapper is best understood by example, which I now show as I wrap HKnotVecter. I begin with Listing 3, which is the file `cpp_HKnotVector.pxd`. In this file I use the `vector` class defined in `libcpp.vector` and include the all the declarations that appear

---

[10]Readers who are unfamiliar with Cython, especially static typing in Cython, are encouraged to read Appendix B before proceeding.

in `HKnotVector.h` (Listing 11 in Appendix A).

Listing 3: `cpp_HKnotVector.pxd`

```
from libcpp.vector cimport vector as cpp_vector

cdef extern from "HKnotVector.h" namespace "hsf":

    cdef cppclass HKnotVector:
        # constructors
        HKnotVector()
        HKnotVector(unsigned int, const cpp_vector[double] &)

        # methods
        unsigned int degree()
        bint isEven()
        bint isOdd()
```

The next part of the wrapper is the Cython interface `HKnotVector.pxd` in Listing 4. This is a very minimal file that declares the `HKnotVector` class and sets up some initial attributes of the class.

Listing 4: `HKnotVector.pxd`

```
from hsfpy cimport cpp_HKnotVector

cdef class HKnotVector:
    cdef void * _inst
    cdef public bint _free_inst
```

The final part and main of the wrapper is `HKnotVector.pyx`, shown here in Listing 5. This is where all attributes and methods declared in either of the interface files are implemented.

Listing 5: `HKnotVector.pyx`

```
cimport numpy as np
from libc.stdlib cimport free
from libcpp.vector cimport vector as cpp_vector
import numpy as np

np.import_array()


cdef class HKnotVector:
    def __cinit__(self, *args, **kwargs):
        self._inst = NULL
        self._free_inst = True

    def __dealloc__(self):
        if self._free_inst:
            free(self._inst)

    # constuctors
    def _constructor1(self):
        self._inst = new cpp_HKnotVector.HKnotVector()
```

```
22   def _constructor2(self, degree, knots):
23       cdef cpp_vector[double] cpp_knots
24       cdef int i
25       cdef int knots_size = len(knots)
26       cpp_knots = cpp_vector[double](<size_t> knots_size)
27       for i in range(knots_size):
28           cpp_knots[i] = <double> knots[i]
29       self._inst = new cpp_HKnotVector.HKnotVector(<unsigned int> long(degree),
                 cpp_knots)
30
31   def __init__(self, *args, **kwargs):
32       if len(args) == 2:
33           self._constructor2(*args, **kwargs)
34       else:
35           self._constructor1(*args, **kwargs)
36
37   # methods
38   def degree(self):
39       cdef unsigned int rtnval
40       rtnval = (<cpp_HKnotVector.HKnotVector *> self._inst).degree()
41       return int(rtnval)
42
43
44   def isEven(self):
45       cdef bint rtnval
46       rtnval = (<cpp_HKnotVector.HKnotVector *> self._inst).isEven()
47       return bool(rtnval)
48
49
50   def isOdd(self):
51       cdef bint rtnval
52       rtnval = (<cpp_HKnotVector.HKnotVector *> self._inst).isOdd()
53       return bool(rtnval)
```

- **Lines 1-6** All necessary items are imported and set up. Notice that neither of the interface files are not actually imported here. If a .pxd and a .pyx file are in the same directory and have the same name, then all things imported or defined in the .pxd are automatically available in the .pyx file.

- **Lines 9-16** Use cdef to declare the class and set up a few special Cython methods. __cinit__ is called immediately after the user tries to create an instance of the class and usually holds the minimial setup required to avoid a SEGFAULT from null pointers. The __dealloc__ method is called when the object is passed through the Python garbage collector and is implemented here to avoid memory leaks.

- **Lines 18-35** Here the overloaded constructor for HKnotVector is set up. Two private methods (private by convention of starting with a single underscore) are implemented to handle each the overloads. The __init__ method is called after __cinit__ when an

HKnotVector instance is created and is implemented to dispatch object creation to one of the overloaded constructors.

- **Lines 37-53** The methods declared in cpp_HKnotVector.pxd are implemented. Pretty much the only thing that needs to happen here is type checking. To do this I use cdef to statically declare variable types and cast objects using < · >.

Now that the wrapper is completed, it needs to be incorporated into a build system and complied into a shared object so that Python can access it. As before, we let Python handle this step using a setup.py file, which I have included in Listing 6. There are only a few differences between this file and the other setup.py files presented earlier. First, in lines 6 and 10 I explicitly specify the include directories for the HKnotVector extension. Also, I setup the hsfpy package with a module named HKnotVector. This happens on lines 8 and 17.

Listing 6: setup.py for Cython wrapper of HKnotVector

```python
from distutils.core import setup
from distutils.extension import Extension
from Cython.Distutils import build_ext
import numpy as np

incdirs = ['..', '.', np.get_include()]

HKnotVector = Extension("hsfpy.HKnotVector",
                        ["hsfpy/HKnotVector.pyx"],
                        include_dirs=incdirs, language="c++")

ext_modules = [HKnotVector]

setup(name='hsfpy',
      cmdclass={'build_ext': build_ext},
      ext_modules=ext_modules,
      packages=['hsfpy']
      )
```

As can be see from this example, wrapping code using Cython provides absolute control over the structure and feel of the wrapper, but it takes more work than, for example, SWIG. I have only wrapped a very small portion of the hsf library in this example, but it illustrates the point. The sheer size of the hsf library makes it unreasonable to construct a wrapper by hand using Cython. Additionally, the core hsf C++ library is still being developed and is therefore liable to change at any time. Trying to keep the Cython wrapper up to date would be a difficult and error-prone task. For these reasons, I decided not to use a by-hand Cython approach in creating hsfpy.

2.4   XDRESS

The final tool I evaluated when creating the Python wrapper for `hsf` is XDress [11]. XDress is a very young project that first appeared on github in April 2013. XDress is written in pure Python and is an automatic Python wrapper generator for C and C++ source. It constructs the wrapper in a three stage process.

1. External (to XDress) parsing tools are run on the source and a static xml representation of the data structures is generated. Currently, XDress uses GCC-XML[12] for C++ parsing and pycparser[13] for C.

2. The generated xml files are parsed and the C-based API is described in terms of an internal XDress typesytem. This typesystem is very dynamic and was designed from the ground up with API generation in mind. It is the main enabling feature of XDress.

3. XDress uses various built-in and/or user-supplied plugins to take the API stored in the typesystem and form Cython bindings.

2.4.1   XDRESSRC.PY

Compared to the other methods discussed here, XDress is very easy to use . The main point of entry for using XDress is to call `xdress` from the command line. When this command is executed (with no extra arguments options) it will scan the current directory for a file named `xdressrc.py`. All the instructions for XDress are put into this single python file. It is easiest to understand the types of instructions that need to be in this file by example, so I present one here.

Listing 7: Sample `xdressrc.py` for HKnotVector

```
1  package = 'package'
2  packagedir = 'output'
3  sourcedir = 'src'
4
5  plugins = ('xdress.stlwrap', 'xdress.autoall', 'xdress.autodescribe',
```

---

[11]XDress is free and open source. The source code is hosted at https://github.com/xdress/xdress and the homepage for the project is http://xdress.org/.

[12]GCC-XML is free and open source. The source code is hosted at https://github.com/gccxml/gccxml and the homepage for the project is http://gccxml.github.io/HTML/Index.html.

[13]pycparser is free and open source. The source code is hosed at https://github.com/eliben/pycparser and the (limited) documentation is found in README.rst in the source.

```
 6              'xdress.cythongen', 'foopack.barplug')
 7
 8  ## Which stl containers we need for this code
 9  stlcontainers = [('vector', 'float64'),
10                  ('set', 'int'),
11                  ('map', 'int', ('map', ('vector', 'uint'), ('set', 'char'))),
12                  ('vector', ('vector', 'float64')),
13                  ('set', 'FooClassBar')
14                  ]
15
16  ## Which classes to create wrappers for.
17  classes = [('FooClass', 'Foo'),
18            ('FooClass', 'Bar', 'Foo', 'FooClassBar'),
19            ]
20
21  functions = [('FooFunc', 'Foo')]
22
23  variables = [('barVar', 'Bar')]
```

- **Lines 1-3** Set the name of the Cython package, the output directory where the Cython wrapper will go, and the name of the directory where the C/C++ source lives.

- **Lines 5-7** This is an optional step where the user can specify which plugins should be run when `xdress` is executed. All but the last plugin (`'foopack.barplug'`) are built-in plugins that come with XDress. They perform the following functions:

  - `'xdress.stlwrap'`: Generates wrapper a for C++ STL objects (see next bullet for more information)

  - `'xdress.autoall'` and `'xdress.autodescribe'`: Parse all included files and enter all objects into the typesystem

  - `'xdress.cythongen'`: Use the generated typesystem to actually write out the Cython files that define the wrapper

  - `'foopack.barplug'`: Run the user supplied plugin `barplug` found in the package `foopack`.

  Note that if the `plugins` list is omitted from this file that XDress will automatically populate this list with the necessary plugins to create the Cython interface.

- **Lines 8-14** Specify which STL containers to create Cython wrappers for. These wrappers will be exposed to Python via custom NumPy dtypes that do all data sharing in memory (no copying, which means a lower memory footprint and faster performance). Notice that the specifications here can take on a nested form to accomodate arbitrary complexity.

Also note that non-native C/C++ types can be specified here, with the restriction that the user-defined types be mentioned in the `classes` list below (see next bullet).

- **Lines 16-19**: An optional list of classes XDress should generate wrappers for. The `classes` object specified here should be a list of tuples. There are various formats for specifying the contents of each tuple, but the format used here is (`'source_name'`, `'source_file'`, `'target_name'`, `'target_file'`). Where `source_name` is the name of the class in C++, `source_file` is the file where the class is defined in C++ and the `target` variants are the name and file for how and where the class should be defined in C++.

- **Lines 21-23**: Optional lists of functions and variables that should be wrapped. The syntax is similar to the syntax for classes.

2.4.2 XDRESS PLUGINS

As can be seen, specifying the API elements that need to be wrapped is straightforward and simple. To complement this simplicity XDress has a very easy to use plugin architecture that gives users absolute control over how the wrapping is handled. The plugin system is not a mere afterthought, but build into the core of how XDress operates. All the major functionality of XDress is modularized into distinct plugins that are executed using this archtecture. This means that user-supplied plugins will be given the same precedence as built-in plugins. To demonstrate some of the possibilities for XDress plugins, I will explain two of the plugins I have written to handle issues encountered with wrapping `hsf`.

The first of these plugins is now a part of XDress and lives in `xdress.descfilter`. This plugin allows the user to instruct XDress to "filter out" certain API components from the generated wrapper. This can be done in one of two ways: 1) specify that functions or methods with certain types in the function signature be excluded or 2) specify that certain methods of a class should be excluded. This flexibility can be useful when there are certain functions that shouldn't be exposed to Python. This is also a useful stop-gap for data types that have not yet been implemented into the XDress type system. In Section 3.1 I will present the actual `xdressrc.py` file currently being developed for `hsf`, which will provide a usage example for `xdress.descfilter`.

The second plugin is also part of XDress and lives in `xdress.doxygen`. This plugin uses dOxygen[14] to output an xml version of in-line documentation contained in the C/C++ source. The plugin then parses this xml, automatically generates Python docstrings, and inserts them into the Cython wrappers. These docstrings have many uses such as to provide information on methods, classes, or functions when these objects are inspected at the Python interpreter, or to be used by a tool like Sphinx[15] in conjunction with other content to produce stylized documentation. In Section 3.2 I will give a preview of what these docstrings look like for `hsfpy` inside of IPython.

Another important item to note is that because I became involved with XDress development at a very early stage and because the `hsf` library utilizes advanced C++ language features, much of the recent and current development of XDress is being driven by the needs that arise in wrapping `hsf`. This, together with the ease and freedom XDress provides, caused me to choose XDress as the tool to use in constructing `hsfpy`. As the `hsf` project moves forward, the relationship with the lead XDress developer, Anthony Scopatz, and the close integration between `hsf` and the XDress development cycle will very helpful and should ensure long-term functionality.

3   RESULTS AND DISCUSSION

XDress is still actively developed and does not yet support all the features of `hsf`. As such, `hsfpy` is still a work on progress. It should be noted, however, that the missing functionality in XDress is minor and should be taken care of before the end of August, 2013. At that time, it will be very easy to finish `hsfpy`. In the remainder of this section I will present the current version of `hsfpy`, including the XDress utilities needed to produce it, as well as show some preliminary usage examples.

---

[14]dOxygen is a common documentation utility for C/C++ projects that gives the user the ability to have specially formatted comments in the source code become stylized documentation elements. dOxygen is free and open source. The code is hosted at `https://github.com/doxygen/doxygen` and the homepage for the project is `http://www.stack.nl/~dimitri/doxygen/`.

[15]Sphinx is a Python package that can automatically create html or pdf (via latex) documentation using the reStructuredText markup language. Additionally, Sphinx can inspect docstrings and turn them into stylized documentation elements, much like dOxygen.

3.1   XDRESS AND HSFPY

All instructions needed to create hsfpy are contained in a single xdressrc.py file, which I give

below in Listing 8. This the actual file used in development as of August 5, 2013. After presenting

the file I will explain key parts.

Listing 8: Actual xdressrc.py for hsfpy

```python
import sys
sys.path.insert(0, '.')

## Do basic setup
package = 'hsfpy' # top-level python package name
packagedir = 'hsfpy' # location of the generated python package
sourcedir = 'src' # location of the original C++ source

# Options for my utils.init_setup plugin
init_filename = '__init__.py'
setup_filename = 'setup.py'
run_setup = False


## List the plugins we need. This step is optional, but we use it because we
# need to filter out some types
plugins = ('xdress.autoall',
           'xdress.autodescribe',
           'xdress.doxygen',
           'xdress.descfilter',
           'xdress.cythongen',
           'xdress.stlwrap',
           'utils.init_setup')

# Which types to ignore or exclude in the wrappers
skiptypes = ['ExtractData', 'HExtractCache', 'HMeshCache', 'istream',
             'basic_istream', 'basic_ostream', 'HTrunkData', 'Activity', 'bool',
             'Cell', 'BFunc', 'HExtract', 'Knot', 'Tree', 'CellFace']

# Which methods to skip in various classes
skipmethods = {
    'HNurbsTree': ['saveFile', 'saveTSplineFile', 'loadStreamBody'],
    'HForest': ['loadFile', 'saveFile', 'loadPointsFile', 'loadStream']
}

## Which stl containers we need for this code
stlcontainers = [
    ('set', 'uint'),
    ('vector', 'float64'), # DoubleVec
    ('vector', 'int32'), # IntVec
    ('vector', ('vector', 'int32')), # IntVecVec
    ('vector', ('vector', 'float64')), # DoubleVecVec
    ('vector', 'HKnotVector'),
    ('vector', ('Point', 3, 'double', False)),
    ('vector', ('Point', 3, 'double', True)),
]

```

```
48 variables = [
49     ('Activity', 'common'),
50     ('MeshType', 'HMesh')
51 ]
52
53 ## Which classes to create wrappers for.
54 classes = [
55     ('HKnotVector', 'HKnotVector'),
56     (('Point', 3, 'double', False), 'Point'),
57     (('Point', 3, 'double', True), 'Point'),
58     ('HNurbs', 'HNurbs'),
59     ('Index', 'common'),
60     ('HNurbsTree', 'HNurbsTree'),
61     ('HForest', 'HForest'),
62     # ('FnHAdapt', 'FnHAdapt'),
63     # ('HExtract', 'HExtract'),
64     # (('HMesh', 'CELL_TYPE'), 'HMesh')
65 ]
66
67 ## Which functions to create wrappers for
68 functions = [
69     ('linearParameterizeNURBS', 'HNurbs')
70 ]
```

- **Lines 1-2** Add the current working directory to the path so the local plugin `utils.init_setup` will run[16].

- **Lines 4-7** Set the package name, package directory, and source directory

- **Lines 9-12** Set options for the `utils.init_setup` plugin written specifically for `hsfpy`. This plugin will take the information contained in `xdressrc.py` and automatically create the necessary `setup.py` file. It will also create an `__init__.py` file in the package directory[17]. Finally, the plugin will automatically run `python setup.py build_ext -inplace` if `run_setup` is set to `True`.

- **Lines 15-23** List the plugins XDress should use in creating `hsfpy`.

- **Lines 25-28** Set the `skiptypes` parameter used by the plugin `xdress.descfilter`. Any time an item in this list appears as the type of an argument or return value of a function or method, that method is skipped when the wrapper is generated. For the most part this is necessary for two reasons: 1) The yet un-implemented features of XDress does not

---

[16]I have included the contents of `utils.init_setup.py` in Appendix C.

[17]The `__init__.py` file is used to manage namespaces. The `utils.init_setup` plugin automatically imports all wrapped functions and classes into the main `hsfpy` namespace.

properly wrap one of these types or 2) There is no need for the Python interface to deal
with these types.

- **Lines 30-34** Set the `skipmethods` parameter used by the `xdress.descfilter` plugin.
  This dict contains class names as keys and a list of method names as values. The method
  names specified in the lists will be excluded from the generated wrapper of the class the
  list pertains to.

- **Lines 36-46** Which C++ standard library types to create NumPy dtypes for. Note the
  inclusion of some non-standard types in lines 43-45.

- **Lines 48-70** Which variables, classes, or functions to create wrappers for. The syntax for
  these lists was explained in Section 2.4.1.

The `setup.py` and `__init__.py` files generated after running xdress are presented in Listing
9 and Listing 10, respectively.

Listing 9: Actual `setup.py` for `hsfpy`

```
1  import os
2  from distutils.core import setup
3  from distutils.extension import Extension
4  from Cython.Distutils import build_ext
5
6  import numpy as np
7
8  incdirs = [os.path.join(os.getcwd(), 'hsfpy'),
9             os.path.join(os.getcwd(), 'src'),
10            np.get_include()]
11
12  # Define extensions
13  xdress_extras = Extension("hsfpy.xdress_extra_types",
14                            ["hsfpy/xdress_extra_types.pyx"],
15                            include_dirs=incdirs, language="c++")
16
17
18  stl_cont = Extension("hsfpy.stlcontainers", ["hsfpy/stlcontainers.pyx"],
19                       include_dirs=incdirs, language="c++")
20
21
22  HKnotVector = Extension("hsfpy.HKnotVector",
23         ["hsfpy/HKnotVector.pyx", "src/HKnotVector.cpp"],
24         include_dirs=incdirs, language="c++")
25
26
27  Point = Extension("hsfpy.Point",
28         ["hsfpy/Point.pyx", "src/Point.cpp"],
29         include_dirs=incdirs, language="c++")
30
```

```
31
32  HNurbs = Extension("hsfpy.HNurbs",
33          ["hsfpy/HNurbs.pyx", "src/HNurbs.cpp"],
34          include_dirs=incdirs, language="c++")
35
36
37  common = Extension("hsfpy.common",
38          ["hsfpy/common.pyx", "src/common.cpp"],
39          include_dirs=incdirs, language="c++")
40
41
42  HNurbsTree = Extension("hsfpy.HNurbsTree",
43          ["hsfpy/HNurbsTree.pyx", "src/HNurbsTree.cpp"],
44          include_dirs=incdirs, language="c++")
45
46
47  HForest = Extension("hsfpy.HForest",
48          ["hsfpy/HForest.pyx", "src/HForest.cpp"],
49          include_dirs=incdirs, language="c++")
50
51
52  ext_modules = [stl_cont,
53          xdress_extras,
54          HKnotVector,
55          Point,
56          HNurbs,
57          common,
58          HNurbsTree,
59          HForest]
60
61  setup(name='hsfpy',
62          cmdclass=dict([('build_ext', build_ext)]),
63          ext_modules=ext_modules,
64          packages=['hsfpy']
65          )
```

Listing 10: Actual `__init__.py` for `hsfpy`

```
1   # import classes
2   from hsfpy.HKnotVector import HKnotVector
3   from hsfpy.Point import Point
4   from hsfpy.HNurbs import HNurbs
5   from hsfpy.common import Index
6   from hsfpy.HNurbsTree import HNurbsTree
7   from hsfpy.HForest import HForest
8
9   # import functions
10  from hsfpy.HNurbs import linearParameterizeNURBS
```

3.2 HSFPY USAGE

Although `hsfpy` is not entirely finished, there is still a usable subset of `hsf` available at the time of writing. The basic data structure in `hsf` is the hierarchical spline forest, which is implemented

as `HForest` and has already been wrapped in `hsfpy`.

4   CONCLUSION

REVIEW: I feel like I am giving a quick summary of everything here. Is that what I am supposed to be doing?

I have presented the basics of the hierarchal spline forests project and its associated software implementation. The `hsf` library is written in C++, which allows the code to be very efficient and generally applicable in many contexts. The long-run goal of `hsf` is to be the most flexible and powerful discretization package available. To broaden the potential user-base, increase the rate of development, and lower the barriers to entry for undergraduate contributions to the project, the focus of this project has been creating a wrapper around `hsf` in a high level language.

I described potential language choices and justified the selection of Python as the high-level target for the wrapper. I evaluated and included functional examples of SWIG, Boost.Python, Cython, and XDress as candidates for creating the Python interface to the library: `hsfpy`. XDress best fit the needs of `hsfpy` and was chosen as the tool for its creation. I then gave a detailed explanation of how XDress was used to create the wrapper, including the contributions I have made to XDress to enable functionality required by `hsfpy`. Finally I presented the current state of the wrapper, showed a usage example, and highlighted various features.

There is still some work to be done before this project is complete. Most of that has to do with extending the capabilities of XDress so that is can handle all the C++ language features `hsf` takes advantage of. After that is complete, XDress will be a very robust tool that will make wrapping future projects and maintaining `hsfpy` quite simple.

REFERENCES

[1] K. Judd, *Numerical Methods in Economics*. MIT Press, 1998, ISBN: 9780262100717. [Online]. Available: http://books.google.com/books/about/Numerical_Methods_in_Economics.html?id=9Wxk_z9HskAC.

[2] T. Hughes, J. Cottrell, and Y. Bazilevs, "Isogeometric analysis: cad, finite elements, nurbs, exact geometry and mesh refinement", *Computer Methods in Applied Mechanics and Engineering*, vol. 194, no. 39-41, 2005, ISSN: 0045-7825. DOI: http://dx.doi.org/10.1016/j.cma.2004.10.008. [Online]. Available: http://www.sciencedirect.com/science/article/pii/S0045782504005171.

[3] D. Schillinger, L. Dede, M. Scott, J. Evans, M. Borden, E. Rank, and T. Hughes, "An isogeometric design-through-analysis methodology based on adaptive hierarchical refinement of nurbs, immersed boundary methods, and t-spline cad surfaces", *The Institute for Computational Engineering and Sciences*, 2013.

A   HSFPY (C++) CODE LISTINGS

Below are the code listings that are used as examples throughout section 2.

Listing 11: Portions of HKnotVector.h

```cpp
#ifndef _H_KNOT_VECTOR_H_
#define _H_KNOT_VECTOR_H_

#include "common.h"
#include <vector>
#include <iostream>

using namespace std;
using namespace util;

namespace hsf
{
  class HKnotVector
  {
    /// A one-dimensional object which stores a knot vector of any degree.
    /// No geometric operations are performed using a knot vector, only basis
    /// function queries. This class is best used in connection with a HNURBS
    ///     object which
    /// stores the geometric information. We do store the extra knot for open
    /// knot vectors. So a degree p knot vector will have p + 1 knots at the
    ///     beginning
    /// and end of the knot vector. We currently don't support periodic knot
    /// vectors although this could be added pretty easily.
    public:

    /// Default constructor
    HKnotVector() : mDeg( 0 ) {}

    /// construct a knot vector from a vector of knots. We assume that p + 1
    ///     repeated
    /// knots exists at the beginning and end of the knot vector.
    HKnotVector( uint degree, const DoubleVec &knots )
      : mDeg( degree ), mKnots( knots )
    {
      getKVecData( mKnots, mGroups, mReverseGroups, mMultipleCount );
    }

    /// A destructor
    ~HKnotVector() {}

    /// Returns the degree of this knot vector.
    uint degree() const { return mDeg; }

    /// Returns true if the knot vector is even.
    bool isEven() const { return degree() % 2 == 0; }

    /// Returns true if the knot vector is odd.
    bool isOdd() const { return !isEven(); }

    protected:

    uint mDeg;
```

```cpp
50 |     DoubleVec mKnots;
51 |     IntVec mGroups;
52 |     IntVecVec mReverseGroups;
53 |     IntVec mMultipleCount;
54 |
55 |     /// Returns group, multiplicity, zcount data for a vector of knots.
56 |     void getKVecData( const DoubleVec &knots, IntVec &knot_groups,
57 |             IntVecVec &reverse_knot_groups, IntVec &multiple_counts ) const
58 |     {
59 |       knot_groups.clear();
60 |       reverse_knot_groups.clear();
61 |       multiple_counts.clear();
62 |       knot_groups.push_back( 0 );
63 |       multiple_counts.push_back( 0 );
64 |       uint group_index = 0;
65 |       uint multiple_count = 0;
66 |       IntVec group;
67 |       group.push_back( 0 );
68 |       for( uint iknot = 1; iknot < knots.size(); ++iknot )
69 |       {
70 |         if( equals( knots[ iknot - 1 ], knots[ iknot ], 1e-8 ) )
71 |         {
72 |           group.push_back( iknot );
73 |           ++multiple_count;
74 |         }
75 |         else
76 |         {
77 |           ++group_index;
78 |           multiple_count = 0;
79 |           reverse_knot_groups.push_back( group );
80 |           group.clear();
81 |           group.push_back( iknot );
82 |         }
83 |         knot_groups.push_back( group_index );
84 |         multiple_counts.push_back( multiple_count );
85 |       }
86 |       reverse_knot_groups.push_back( group );
87 |     }
88 |   };
89 | }
90 | #endif
```

Listing 12: Portions of common.h

```cpp
 1 | #ifndef _UTIL_COMMON_H_
 2 | #define _UTIL_COMMON_H_
 3 |
 4 | #include <climits>
 5 | #include <iostream>
 6 | #include <vector>
 7 | #include <set>
 8 | #include <map>
 9 | #include <cmath>
10 | #include <string>
11 | #include <assert.h>
12 |
13 | /// common definitions needed throughout the hsf library
14 |
```

```cpp
15 typedef unsigned int uint;
16 typedef unsigned long ulong;
17 typedef unsigned short ushort;
18 typedef unsigned char uchar;
19
20 using namespace std;
21
22 namespace util
23 {
24   /// Clamps the values to a determined range. The values
25   /// 'minimum' and 'maxmimum' must be of a type that can be
26   /// cast to the same type as 'value', and must be less-than
27   /// comparable with value's type as well.
28   template< typename T, typename T2, typename T3 >
29   inline T numClamp( T value, T2 minimum, T3 maximum )
30   {
31     if( value < minimum )
32       return minimum;
33     if( maximum < value )
34       return maximum;
35     return value;
36   }
37
38   /// This form is a little inconvenient, but is the basis of most other
39   /// ways of measuring equality.
40   inline bool equals( double a, double b, double tolerance )
41   {
42     // This method has been benchmarked, and it's pretty fast.
43     return ( a == b ) ||
44       ( ( a <= ( b + tolerance ) ) &&
45       ( a >= ( b - tolerance ) ) );
46   }
47
48   typedef std::vector< double > DoubleVec;
49   typedef std::vector< int > IntVec;
50   typedef std::vector< IntVec > IntVecVec;
51 }
52 #endif
```

B CYTHON TYPE EXAMPLE

The main point of entry for adding static types in Cython is the `cdef` keyword. This can be used before any object to assign a type to it. All C types can be used as valid `cdef` declarations: numeric types, structs, unions, pointers, ect. In addition, many Python types like `list` or `dict` have been optimized to get performance gains when `cdef` is used to declare their type. When using `cdef`, Cython will generate C code that does automatic type conversion between related Python and C types. The end result of code that has been properly typed using `cdef` is much

faster code - sometimes faster by orders of magnitude.

To demonstrate the use of the `cdef` keyword I will show Python and Cython versions of a pairwise-distance function. This function takes in an $n \times m$ matrix that represents n points in m dimensions and it will return an $n \times n$ matrix containing the Euclidean distance between each point in the input array and every other point in that array. I show Python and Cython versions below and then explain the differences:

Listing 13: `pairs.py`: Pure Python pairwise distance function

```python
from math import sqrt
import numpy as np


def dist(x):
    n = x.shape[0]
    m = x.shape[1]
    ret = np.empty((n, n))
    for i in range(n):
        for j in range(n):
            d = 0.0
            for k in range(m):
                tmp = x[i, k] - x[j, k]
                d += tmp * tmp
            ret[i, j] = sqrt(d)
    return ret
```

Listing 14: `cy_pairs.pyx`: Cython pairwise distance function

```python
from libc.math cimport sqrt
import numpy as np


cpdef dist(double[:, ::1] x):
    cdef int n = x.shape[0]
    cdef int m = x.shape[1]
    cdef double[:, ::1] ret = np.empty((n, n))
    cdef double d, tmp
    cdef int i, j, k
    for i in range(n):
        for j in range(n):
            d = 0.0
            for k in range(m):
                tmp = x[i, k] - x[j, k]
                d += tmp * tmp
            ret[i, j] = sqrt(d)
    return ret
```

- **Line 1** Cython exposes the much of the C standard library via `libc.<headerName>`. The `sqrt` function from the standard library is a bit faster than the one from Python's built-in

math package. Note that I must use the Cython keyword `cimport` to access this function.

- **Line 5** Notice the use of the keyword `cpdef`. This keyword is used to define functions or classes that need to be callable from both Python and C. Were I to have used `cdef` here, the function would be translated to a C function and I would not be able to call it from Python. Behind the scenes `cpdef` instructs the Cython to C translator to make two versions of the function: one for Python use and the other for C use.

- **Line 5** Also note that on line 5 I declare a Cython typed memoryview using `double[:, ::1]`. This statement tells Cython that `x` will be a two dimensional array of doubles. In addition, the `::1` in the second position tells Cython that `x` will be C-contiguous[18]. This allows the generated C code to use natural C array operations on `x`.

- **Lines 6-10** Here I give static types to all variables local to the function. Note the use of the typed memory view again on line 8. Also note that Cython requires types to be declared at the top level of a function. For that reason, I declared `d` and `tmp` as double and `i, j, k` as int before entering first `for` loop.

- The rest of the function is identical to the pure Python version.

In order to use the Cython version of the function, we must instruct Cython to translate it to C and then compile it for Python use. There are many ways to do this, but as with SWIG it is easiest to let Python handle it for us using a `setup.py` file. A setup.py file for this function appears below:

Listing 15: `setup.py` file for Cython pairwise distance

```
1  from distutils.core import setup
2  from Cython.Build import cythonize
3
4  setup(name="Pairwise distance", ext_modules=cythonize('cy_pairs.pyx'))
```

This file is very simple: lines 1 and 2 import the `setup` and `cythonize` functions and line 4 calls the `setup` function where the extension modules are given using the `cythonize` function. The only remaining step is to build the extension using the command used to build the SWIG extension above. I repeat the command here:

---

[18]Note that by default all NumPy arrays are C-contiguous.

```
        python setup.py build_ext --inplace
```

I timed both of these functions using x = np.random.randn(1500, 5) as the input array. Both functions returned the exact same answer, but the execution time was very different. The Python function took 21.2 seconds to execute, whereas the Cython version only took 75.6 milliseconds: a speedup of over $280x$ [19]!

C INIT_SETUP.PY XDRESS PLUGIN

The contents of the utils.init_setup plugin mentioned in Section 3.1 appear below.

Listing 16: XDress plugin init_setup.py

```
1  import sys
2  import os
3  import io
4  from subprocess import call
5  from xdress.plugins import Plugin
6
7  if sys.version < 3:
8      basestring = str
9
10 import collections
11
12
13 class OrderedSet(collections.MutableSet):
14
15     def __init__(self, iterable=None):
16         self.end = end = []
17         end += [None, end, end]      # sentinel node for doubly linked list
18         self.map = {}                # key --> [key, prev, next]
19         if iterable is not None:
20             self |= iterable
21
22     def __len__(self):
23         return len(self.map)
24
25     def __contains__(self, key):
26         return key in self.map
27
28     def add(self, key):
29         if key not in self.map:
30             end = self.end
31             curr = end[1]
32             curr[2] = end[1] = self.map[key] = [key, curr, end]
33
```

---

[19]I also have a more optimized version of the Cython code that only takes 14.9 milliseconds to run. While that shows a speed improvement of over 1400x, it makes use of some advanced Cython features that are beyond the scope of this report.

```python
34      def discard(self, key):
35          if key in self.map:
36              key, prev, next = self.map.pop(key)
37              prev[2] = next
38              next[1] = prev
39
40      def __iter__(self):
41          end = self.end
42          curr = end[2]
43          while curr is not end:
44              yield curr[0]
45              curr = curr[2]
46
47      def __reversed__(self):
48          end = self.end
49          curr = end[1]
50          while curr is not end:
51              yield curr[0]
52              curr = curr[1]
53
54      def pop(self, last=True):
55          if not self:
56              raise KeyError('set is empty')
57          key = self.end[1][0] if last else self.end[2][0]
58          self.discard(key)
59          return key
60
61      def __repr__(self):
62          if not self:
63              return '%s()' % (self.__class__.__name__,)
64          return '%s(%r)' % (self.__class__.__name__, list(self))
65
66      def __eq__(self, other):
67          if isinstance(other, OrderedSet):
68              return len(self) == len(other) and list(self) == list(other)
69          return set(self) == set(other)

_extension = """\
{tarfile} = Extension("{pack}.{tarfile}",
\t["{packdir}/{tarfile}.pyx", "{srcdir}/{srcfile}.cpp"],
\tinclude_dirs=incdirs, language="c++")
"""

_stl_extention = """\
stl_cont = Extension("{pack}.stlcontainers", ["{pack}/stlcontainers.pyx"],
                     include_dirs=incdirs, language="c++")
"""
_xd_extras_ext = """\
xdress_extras = Extension("{pack}.xdress_extra_types",
                          ["{pack}/xdress_extra_types.pyx"],
                          include_dirs=incdirs, language="c++")
"""

_setup_main = """\
import os
from distutils.core import setup
from distutils.extension import Extension
from Cython.Distutils import build_ext
```

```
 92 
 93 import numpy as np
 94 
 95 incdirs = [os.path.join(os.getcwd(), '{packdir}'),
 96            os.path.join(os.getcwd(), '{srcdir}'),
 97            np.get_include()]
 98 
 99 # Define extensions
100 {extensions}
101 
102 ext_modules = {ext_mod_list}
103 
104 setup(name='{pack}',
105       cmdclass=dict([('build_ext', build_ext)]),
106       ext_modules=ext_modules,
107       packages=['{pack}']
108       )
109 """
110 
111 _import_obj = "from {pack}.{tarfile} import {obj}\n"
112 
113 
114 _init_main = """\
115 # import classes
116 {classes}
117 # import functions
118 {functions}
119 
120 """
121 
122 
123 class XDressPlugin(Plugin):
124     """Plugin for generating __init__.py and setup.py."""
125     defaultrc = {'init_filename': '__init__.py',
126                  'setup_filename': 'setup.py',
127                  'run_setup': False}
128 
129     def execute(self, rc):
130         package = rc.package
131         packagedir = rc.packagedir
132         srcdir = rc.sourcedir
133 
134         classes = rc.classes
135         functions = rc.functions
136 
137         classes_set = OrderedSet()
138         ext_set = OrderedSet()
139         str_funcs = ''
140         e_mod_list = OrderedSet()
141 
142         if os.path.isfile(packagedir + os.path.sep + 'xdress_extra_types.pyx'):
143             ext_set.add(_xd_extras_ext.format(pack=package))
144             e_mod_list.add('stl_cont')
145 
146         if os.path.isfile(packagedir + os.path.sep + 'stlcontainers.pyx'):
147             ext_set.add(_stl_extention.format(pack=package))
148             e_mod_list.add('xdress_extras')
149 
150         for cc in classes:
151             srcfile = cc.srcfile
```

33

```
152     tarfile = cc.tarfile
153     c_tarname = cc.tarname
154     if isinstance(c_tarname, basestring):
155         tarname = c_tarname
156     else:
157         tarname = c_tarname[0]
158
159     classes_set.add(_import_obj.format(pack=package,
160                                        tarfile=tarfile,
161                                        obj=tarname))
162
163     ext_set.add(_extension.format(pack=package,
164                                   tarfile=tarfile,
165                                   packdir=packagedir,
166                                   srcdir=srcdir,
167                                   srcfile=srcfile
168                                   ))
169     e_mod_list.add(tarfile)
170
171 for ff in functions:
172     srcfile = ff.srcfile
173     tarfile = ff.tarfile
174     c_tarname = ff.tarname
175     if isinstance(c_tarname, basestring):
176         tarname = c_tarname
177     else:
178         tarname = c_tarname[0]
179
180     str_funcs += _import_obj.format(pack=package,
181                                     tarfile=tarfile,
182                                     obj=tarname)
183
184     ext_set.add(_extension.format(pack=package,
185                                   tarfile=tarfile,
186                                   packdir=packagedir,
187                                   srcdir=srcdir,
188                                   srcfile=srcfile
189                                   ))
190     e_mod_list.add(tarfile)
191
192 # Create the string for the extension modules
193 str_e_mod_list = '['
194 str_e_mod_list += ',\n\t'.join(e_mod_list)
195 str_e_mod_list += ']'
196
197 # Create the string for the classes
198 str_classes = ''.join(classes_set)
199
200 # Create string for the extensions
201 str_extensions = '\n\n'.join(ext_set)
202
203 # Write the __init__.py
204 init_name = packagedir + os.path.sep + rc.init_filename
205 print('init_setup: Writing %s' % init_name)
206 init_txt = _init_main.format(classes=str_classes, functions=str_funcs)
207 with io.open(init_name, 'wb') as f:
208     f.write(init_txt)
209
210 setup_txt = _setup_main.format(extensions=str_extensions,
```

```
211                                              ext_mod_list=str_e_mod_list,
212                                              pack=package,
213                                              packdir=packagedir,
214                                              srcdir=srcdir)
215
216         # Write the setup.py
217         print('init_setup: Writing %s' % rc.setup_filename)
218         with io.open(rc.setup_filename, 'wb') as f:
219             f.write(setup_txt)
220
221         if rc.run_setup:
222             print('#' * 72 + '\nRunning setup.py')
223             call(['python', 'setup.py', 'build_ext', '--inplace'])
```